

Name: Prajakta Mahangare

Roll No: 34

PRN NO.: 12520047

Class: CSE(AI)-SEDA Batch: B2

Assignment No.6 Use the diabetes data set from UCI and Pima Indians Diabetes data set for performing the following: a. Univariate analysis: Frequency, Mean, Median, Mode, Variance, Standard Deviation, Skewness and Kurtosis b. Bivariate analysis: Linear and logistic regression modeling c. Multiple Regression analysis d. Also compare the results of the above analysis for the two data sets.

Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score
from scipy.stats import skew, kurtosis
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, classification_report
```

Load Datasets

```
df_pima = pd.read_csv("diabetes.csv")
print("Pima Dataset Shape:", df_pima.shape)
df_pima.head()
```

Pima Dataset Shape: (768, 9)

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Next steps: [Generate code with df_pima](#) [New interactive sheet](#)

```
diabetes = load_diabetes()
df_uci = pd.DataFrame(diabetes.data, columns=diabetes.feature_names)
df_uci['target'] = diabetes.target
print("UCI Dataset Shape:", df_uci.shape)
df_uci.head()
```

UCI Dataset Shape: (442, 11)

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6	target
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	-0.002592	0.019907	-0.017646	151.0
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	-0.039493	-0.068332	-0.092204	75.0
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	-0.002592	0.002861	-0.025930	141.0
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	0.034309	0.022688	-0.009362	206.0
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	-0.002592	-0.031988	-0.046641	135.0

Next steps: [Generate code with df_uci](#) [New interactive sheet](#)

Data CLeaning

```
print("Pima Missing Values:\n", df_pima.isnull().sum())
print("UCI Missing Values:\n", df_uci.isnull().sum())
```

```
Pima Missing Values:
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
UCI Missing Values:
age      0
sex      0
bmi      0
bp       0
s1       0
s2       0
s3       0
s4       0
s5       0
s6       0
target   0
dtype: int64
```

Replace Invalid Zeros (Pima Dataset)

```
cols = ['Glucose','BloodPressure','SkinThickness','Insulin','BMI']
df_pima[cols] = df_pima[cols].replace(0, np.nan)

df_pima.fillna(df_pima.mean(), inplace=True)
df_pima.drop_duplicates(inplace=True)
```

Univariate Analysis

Frequency

```
print("Pima Outcome Frequency:\n", df_pima['Outcome'].value_counts())
print("\nUCI Target Frequency (Top 10):\n", df_uci['target'].value_counts().head(10))
```

```
Pima Outcome Frequency:
Outcome
0      500
1      268
Name: count, dtype: int64

UCI Target Frequency (Top 10):
target
200.0    6
72.0     6
178.0    5
71.0     5
90.0     5
97.0     4
55.0     4
53.0     4
131.0    4
150.0    4
Name: count, dtype: int64
```

Mean, Median, Mode

```
print("Pima Mean:\n", df_pima.mean())
print("UCI Mean:\n", df_uci.mean())

print("Pima Median:\n", df_pima.median())
print("UCI Median:\n", df_uci.median())

print("Pima Mode:\n", df_pima.mode().iloc[0])
print("UCI Mode:\n", df_uci.mode().iloc[0])
```

```
s3      -6.028360e-18
s4      -1.788100e-17
s5       9.243486e-17
s6       1.351770e-17
target   1.521335e+02
dtype: float64
Pima Median:
Pregnancies      3.000000
Glucose          117.000000
BloodPressure    72.202592
SkinThickness    29.153420
Insulin          155.548223
BMI              32.400000
DiabetesPedigreeFunction  0.372500
Age              29.000000
Outcome          0.000000
dtype: float64
UCI Median:
age      0.005383
sex      -0.044642
bmi      -0.007284
bp       -0.005670
s1       -0.004321
s2       -0.003819
s3       -0.006584
s4       -0.002592
s5       -0.001947
s6       -0.001078
target   140.500000
dtype: float64
Pima Mode:
Pregnancies      1.000000
Glucose          99.000000
BloodPressure    70.000000
SkinThickness    29.153420
Insulin          155.548223
BMI              32.000000
DiabetesPedigreeFunction  0.254000
Age              22.000000
Outcome          0.000000
Name: 0, dtype: float64
UCI Mode:
age      0.016281
sex      -0.044642
bmi      -0.030996
bp       -0.040099
s1       -0.037344
s2       -0.001001
s3       -0.013948
s4       -0.039493
s5       -0.018114
s6       0.003064
target   72.000000
Name: 0, dtype: float64
```

Variance & Standard Deviation

```
print("Pima Variance:\n", df_pima.var())
print("UCI Variance:\n", df_uci.var())

print("Pima Std Dev:\n", df_pima.std())
print("UCI Std Dev:\n", df_uci.std())
```

```
Pima Variance:
Pregnancies      11.354056
Glucose          926.346983
BloodPressure    146.321591
SkinThickness    77.280660
Insulin          7228.588766
BMI              47.267706
DiabetesPedigreeFunction  0.109779
Age              138.303046
Outcome          0.227483
dtype: float64
UCI Variance:
age      0.002268
sex      0.002268
bmi      0.002268
bp       0.002268
s1       0.002268
s2       0.002268
s3       0.002268
s4       0.002268
s5       0.002268
s6       0.002268
target   5943.331348
```

```
dtype: float64
Pima Std Dev:
  Pregnancies      3.369578
  Glucose          30.435949
  BloodPressure    12.096346
  SkinThickness    8.790942
  Insulin          85.021108
  BMI              6.875151
  DiabetesPedigreeFunction  0.331329
  Age             11.760232
  Outcome          0.476951
dtype: float64
UCI Std Dev:
  age      0.047619
  sex      0.047619
  bmi      0.047619
  bp       0.047619
  s1       0.047619
  s2       0.047619
  s3       0.047619
  s4       0.047619
  s5       0.047619
  s6       0.047619
  target   77.093005
dtype: float64
```

Skewness & Kurtosis

```
print("PIMA SKEWNESS & KURTOSIS")
for col in df_pima.columns:
    print(col, "Skew:", skew(df_pima[col]), "Kurtosis:", kurtosis(df_pima[col]))

print("\nUCI SKEWNESS & KURTOSIS")
for col in df_uci.columns:
    print(col, "Skew:", skew(df_uci[col]), "Kurtosis:", kurtosis(df_uci[col]))
```


```
PIMA SKEWNESS & KURTOSIS
Pregnancies Skew: 0.8999119408414357 Kurtosis: 0.15038273760948462
Glucose Skew: 0.531677628850459 Kurtosis: -0.2652767871408801
BloodPressure Skew: 0.1370370472689305 Kurtosis: 1.082847367848931
SkinThickness Skew: 0.8205664594675018 Kurtosis: 5.3718544792635345
Insulin Skew: 3.013183811523944 Kurtosis: 15.078748264774752
BMI Skew: 0.5970835583878779 Kurtosis: 0.9057125569396862
DiabetesPedigreeFunction Skew: 1.9161592037386292 Kurtosis: 5.550792047551203
Age Skew: 1.127389259531697 Kurtosis: 0.6311769413798585
Outcome Skew: 0.6337757030614577 Kurtosis: -1.5983283582089547

UCI SKEWNESS & KURTOSIS
age Skew: -0.23059556012398927 Kurtosis: -0.6771986900987508
sex Skew: 0.1269518188640044 Kurtosis: -1.983883235687121
bmi Skew: 0.5961166556214368 Kurtosis: 0.08047812866813064
bp Skew: 0.28967103827590446 Kurtosis: -0.5403332293804954
s1 Skew: 0.3768238227280956 Kurtosis: 0.2167770625244283
s2 Skew: 0.4351087583570202 Kurtosis: 0.5810556912433786
s3 Skew: 0.7965401530972942 Kurtosis: 0.9568955290670442
s4 Skew: 0.7328756796573327 Kurtosis: 0.42584638305986955
s5 Skew: 0.2907626799653419 Kurtosis: -0.14639566230288814
s6 Skew: 0.20721035056151146 Kurtosis: 0.22070114138631514
target Skew: 0.43906639932477265 Kurtosis: -0.8866436055681386
```

VISUALIZATION

```
plt.figure(figsize=(5,4))
df_pima['Glucose'].hist()
plt.title("Glucose Distribution (Pima)")
plt.tight_layout()
plt.show()



plt.figure(figsize=(5,4))
df_uci['bmi'].hist()
plt.title("BMI Distribution (UCI)")
plt.tight_layout()
plt.show()
```



Correlation Heatmap

```
plt.figure(figsize=(10,6))
sns.heatmap(df_pima.corr(), annot=True, cmap='coolwarm')
plt.title("Pima Dataset Correlation Heatmap")
plt.show()

plt.figure(figsize=(10,6))
sns.heatmap(df_uci.corr(), annot=True, cmap='coolwarm')
plt.title("UCI Dataset Correlation Heatmap")
plt.show()
```



Bivariate Analysis

Linear Regression(UCI Dataset)

```
X_uci = df_uci.drop('target', axis=1)
y_uci = df_uci['target']

X_train_u, X_test_u, y_train_u, y_test_u = train_test_split(
    X_uci, y_uci, test_size=0.2, random_state=42)

lin_model = LinearRegression()
lin_model.fit(X_train_u, y_train_u)

y_pred_u = lin_model.predict(X_test_u)

print("UCI Linear Regression")
print("MSE:", mean_squared_error(y_test_u, y_pred_u))
print("R2 Score:", r2_score(y_test_u, y_pred_u))
```

UCI Linear Regression
MSE: 2900.193628493482
R2 Score: 0.4526027629719195

Logistic Regression (Pima Dataset)

```
X_pima = df_pima.drop('Outcome', axis=1)
y_pima = df_pima['Outcome']

X_train_p, X_test_p, y_train_p, y_test_p = train_test_split(
    X_pima, y_pima, test_size=0.2, random_state=42)
```



```
log_model = LogisticRegression(max_iter=2000)
log_model.fit(X_train_p, y_train_p)

y_pred_p = log_model.predict(X_test_p)

print("Pima Logistic Regression")
print("Accuracy:", accuracy_score(y_test_p, y_pred_p))
```

Pima Logistic Regression
Accuracy: 0.7532467532467533

MULTIPLE REGRESSION ANALYSIS(UCI)

```
multi_model = LinearRegression()
multi_model.fit(X_train_u, y_train_u)

print("Multiple Regression R2 Score:",
      multi_model.score(X_test_u, y_test_u))
```

Multiple Regression R2 Score: 0.4526027629719195


Confusion Matrix

```
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test_p, y_pred_p)

plt.figure(figsize=(4,3))
sns.heatmap(cm,
            annot=True,
            fmt='d',
            cmap='Blues',
            xticklabels=['No Diabetes', 'Diabetes'],
            yticklabels=['No Diabetes', 'Diabetes'])

plt.title("Confusion Matrix - Pima Dataset")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



Precision, Recall, F1

```
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report

print("Precision:", precision_score(y_test_p, y_pred_p))
print("Recall:", recall_score(y_test_p, y_pred_p))
print("F1 Score:", f1_score(y_test_p, y_pred_p))

print("\nClassification Report:\n",
      classification_report(y_test_p, y_pred_p))
```

Precision: 0.6666666666666666
Recall: 0.6181818181818182
F1 Score: 0.6415094339622641

Classification Report:
precision recall f1-score support


0	0.80	0.83	0.81	99
1	0.67	0.62	0.64	55
accuracy			0.75	154
macro avg	0.73	0.72	0.73	154
weighted avg	0.75	0.75	0.75	154

```
pima_mean = df_pima.mean().mean()
uci_mean = df_uci.mean().mean()

plt.figure(figsize=(4,3))
plt.bar(['Pima Dataset','UCI Dataset'],
        [pima_mean, uci_mean])
plt.title("Average Mean Comparison")
plt.ylabel("Average Mean")
plt.show()

uci_r2 = r2_score(y_test_u, y_pred_u)
pima_accuracy = accuracy_score(y_test_p, y_pred_p)

plt.figure(figsize=(4,3))
plt.bar(['UCI (R2 Score)','Pima (Accuracy)'],
        [uci_r2, pima_accuracy])
plt.title("Model Performance Comparison")
plt.ylabel("Score")
plt.show()
```



```
plt.figure(figsize=(5,4))
plt.bar(['Pima','UCI'], [pima_mean, uci_mean])
plt.title("Mean Comparison")
plt.tight_layout()
plt.show()
```

